

Recommender Systems Workshop

Dr. Barry Shepherd
Institute of Systems Science
National University of Singapore
Email: barryshepherd@nus.edu.sg



© 2017 NUS. The contents contained in this document may not be reproduced in any form or by any means, without the written permission of ISS, NUS, other than for the purpose for which it has been supplied.

Workshop Synopsis

We use the *movielens* database – this contains the ratings of 943 users on 1682 movies - the goal is to make movie recommendations to a test set of users.

1. Hand-code and test a simple **user-based** collaborative filtering recommender system using R (use my *starter* code if you wish)
2. Repeat using **alternative similarity measures**, compare results
3. Convert your system to **item-based** collaborative filtering , compare results with step1 & 2



The MovieLens Data Set

- Each user has rated at least 20 movies
- Each ratings record has the format: UserID, MovieID, Rating, Timestamp
 - The data is randomly ordered. Users and items are numbered consecutively from 1.
 - Ratings are made on a 5-star scale (whole-star ratings only)
 - Timestamp is represented in seconds since 1/1/1970 UTC

UserID	movie	rating	datetime
1	61	4	878542420
1	189	3	888732928
1	33	4	878542699
1	160	4	875072547
1	20	4	887431883
1	202	5	875072442
1	171	5	889751711
1	265	4	878542441

1	Toy Story (1995)
2	GoldenEye (1995)
3	Four Rooms (1995)
4	Get Shorty (1995)
5	Copcat (1995)
6	Shanghai Triad (Yao a yao y)
7	Twelve Monkeys (1995)
8	Babe (1995)
9	Dead Man Walking (1995)
10	Richard III (1995)

Get movie names from a separate file

MovieLens Dataset in Tabular Format

- I have converted the movie lens data set into tabular format*

User	movie1	movie2	movie3	movie4	movie5	movie6	etc....
1	2	5	4		3	1	
2		3		5	3	1	
3			2	3			

- *Note that a much bigger dataset is also available containing 1,000,209 anonymous ratings of approximately 3,900 movies made by 6,040 MovieLens users who joined MovieLens in 2000.
 - Ambitious students may use this if you wish
 - This has not been converted to tabular format

The Basic User-User CF algorithm

- Each user is represented by a single record (vector) containing a set of properties (features) – these typically the ratings or purchases of some of the items to be recommended (e.g. movies)
- To make a recommendation to a user
 - Compute the similarity of that user to all other users in the database (typically we use the Pearson coefficient)
 - For every item NOT rated or bought by the user
 - Compute the weighted average rating of all the other users for that item (or just consider the K nearest neighbours)
 - Weighted average = $(\sum_{\text{users}} \text{Item Rating} * \text{User Similarity}) / \sum_{\text{users}} \text{User Similarity}$
 - Recommend the item with the biggest weighted average rating


A simple example using R

- First load the data* into a data frame

```
> users <- read.csv("simplemovies.csv")
> users
```

	user	LadyInWater	SnakesOnPlane	JustMyLuck	Superman	Dupree	NightListener
1	Rose	2.5	3.5	3.0	3.5	2.5	3.0
2	Seymour	3.0	3.5	1.5	5.0	3.0	3.5
3	Philips	2.5	3.0	NA	3.5	NA	4.0
4	Puig	NA	3.5	3.0	4.0	2.5	4.5
5	LaSalle	3.0	4.0	2.0	3.0	2.0	3.0
6	Matthews	3.0	4.0	NA	5.0	3.5	3.0
7	Toby	NA	4.5	NA	4.0	1.0	NA

*Example data taken from:
"Collective Intelligence", O'Reilly



A simple example with sample R code

- To compute the similarity between users we use $\text{cor}(X,Y)$. This computes the Pearson correlation between the columns of X and Y. Hence we first need to transpose the data using $t(..)$ so that users are on the columns

```
> items <- as.data.frame(t(users[,2:ncol(users)]))
> colnames(items) <- users[,1]
> items
```

	Rose	Seymour	Philips	Puig	LaSalle	Matthews	Toby
LadyInWater	2.5	3.0	2.5	NA	3	3.0	NA
SnakesOnPlane	3.5	3.5	3.0	3.5	4	4.0	4.5
JustMyLuck	3.0	1.5	NA	3.0	2	NA	NA
Superman	3.5	5.0	3.5	4.0	3	5.0	4.0
Dupree	2.5	3.0	NA	2.5	2	3.5	1.0
NightListener	3.0	3.5	4.0	4.5	3	3.0	NA

A simple example with sample R code

- Set $\text{use} = \text{"pairwise.complete.obs"}$ to ignore missing values during the correlation computation otherwise most correlations are NA (uncomputible)

```
> cor(items,items)
```

	Rose	Seymour	Philips	Puig	LaSalle	Matthews	Toby
Rose	1.0000000	0.4950738	NA	NA	0.5940885	NA	NA
Seymour	0.4950738	1.0000000	NA	NA	0.5294118	NA	NA
Philips	NA	NA	NA	NA	NA	NA	NA
Puig	NA	NA	NA	NA	NA	NA	NA
LaSalle	0.5940885	0.5294118	NA	NA	1.0000000	NA	NA
Matthews	NA	NA	NA	NA	NA	NA	NA
Toby	NA	NA	NA	NA	NA	NA	NA

```
>
> cor(items, items, use="pairwise.complete.obs")
```

	Rose	Seymour	Philips	Puig	LaSalle	Matthews	Toby
Rose	1.0000000	0.4950738	0.4045199	0.56694671	0.5940885	0.74701788	0.9912407
Seymour	0.4950738	1.0000000	0.4472136	0.56694671	0.5294118	0.87287156	0.5921369
Philips	0.4045199	0.4472136	1.0000000	1.00000000	-0.2581989	0.13483997	-1.0000000
Puig	0.5669467	0.5669467	1.0000000	1.00000000	0.5669467	0.02857143	0.8934051
LaSalle	0.5940885	0.5294118	-0.2581989	0.56694671	1.0000000	0.21128856	0.9244735
Matthews	0.7470179	0.8728716	0.1348400	0.02857143	0.2112886	1.00000000	0.6628490
Toby	0.9912407	0.5921369	-1.0000000	0.89340515	0.9244735	0.66284898	1.0000000

A simple example with sample R code

- To obtain Recommendations for Toby...

```
#(1) get similarity of Toby to all other users

#method1
tobycol = grep("Toby", colnames(items))
sims = cor(items[,tobycol], items[,-tobycol],use="pairwise.complete.obs")

#method2
sims = cor(items[, "Toby"], items[,!names(items) %in% c("Toby")],use="pairwise.complete.obs")
sims = sims[,!is.na(sims)] # some users may have no ratings at all or none in common with Toby

#(2) for each movie (row), compute weighted average rating for all other users except Toby's
wavrats = apply(items[,names(sims)],1,function(x) weighted.mean(x, sims+1, na.rm=TRUE))
wavrats = wavrats[!is.na(wavrats)] #some movies may have no ratings at all (wavrat will be NA)

#(3) eliminate the movies already seen by Toby
|
notseenitems = row.names(items[is.na(items[, "Toby"]),])
predrats = wavrats[notseenitems]

#(4) sort in descending order of predicted rating and select the top 5
sort(predrats[!is.na(predrats)], decreasing = TRUE)[1:min(5,length(predrats))]
```

A simple example with sample R code

- Generalising the previous slide's code into a function:

```
getrecommendations <- function(target) {

  # compute similarity between targetuser and all other users
  sims <- cor(items[,target],items[,!names(items) %in% c(target)],use="pairwise.complete.obs")
  sims <- sims[1,!is.na(sims)]

  # for each item compute weighted average of all the other user ratings
  wavrats = apply(items[,names(sims)],1,function(x) weighted.mean(x, sims+1, na.rm=TRUE))
  wavrats = wavrats[!is.na(wavrats)]

  # remove items already rated by the user
  notseenitems = row.names(items[is.na(items[,target]),])
  t = wavrats[notseenitems]
  sort(t[!is.na(t)] , decreasing = TRUE)[1:min(5,length(t))] # get top 5 items
}

> getrecommendations("Toby")
NightListener LadyInWater JustMyLuck
3.401162 2.861154 2.417304
> .
```

The *apply()* function applies a function to every row (or column) in a data frame. Make sure you read the manual to understand fully how it works!

We add +1 to *sims* to ensure they are >=0

Testing the Recommendations (1)

- Split the available data into training and test sets
- Consider each test user in turn:

User	movie1	movie2	movie3	movie4	movie5	movie6	etc....
Test user	2	5	4		3	1	

↑ ↑ ↑ ↑ ↑

For each movie rated by a test user:

- set the movie rating to blank (NA) – but keep a copy
- make a prediction for that movie using the training data
- compare the prediction with actual rating:
 $error = abs(predicted\ rating - actual\ rating)$
- keep a running total of the errors & number of tests:
 $totalerror = totalerror + error$
 $cnt = cnt + 1$
- restore the blank movie rating

Do this for all test users

At the end compute the overall MAE (mean average error)

$MAE = totalerror / cnt$

Example Testing Code

```
testusernames = sample(names(items), 2) # identify 2 user randomly for testing
trainusernames = setdiff(names(items), testusernames) # take remaining users for training

#test recommendations for all users
testall <- function() {
  totterr = 0
  for (user in testusernames) {
    mae = testuser(user)
    cat("mae for ", user, " is ", mae, "\n");
    totterr = totterr + mae
  }
  cat(sprintf("AVERAGE MAE=%0.4f\n", totterr/length(testusernames)))
}

#test recommendations for one user
testuser <- function(target) {
  testitems = row.names(items[!is.na(items[,target]),])
  targetdata = items[testitems,target]
  names(targetdata) = testitems
  traindata = items[testitems,trainusernames]
  totterr = valid = 0
  for (item in testitems) {
    truerating = targetdata[item]
    targetdata[item] = NA
    sims = cor(targetdata,traindata,use="pairwise.complete.obs")
    sims = sims[,!is.na(sims)]
    prediction = weighted.mean(traindata[item,names(sims)], sims+1, na.rm=TRUE)
    if (!is.na(prediction)) {
      totterr = totterr + abs(prediction - unname(truerating))
      valid = valid + 1
    }
    targetdata[item] = truerating
  }
  return(totterr/valid)
}
```

Testing the Recommendations (2)

- What does a MAE of (say) 1.19 mean in practice? Is it good or bad?
- We need to know how many predictions would actually be made and how many would likely be received favorably by the user?
- To answer this we need a Confusion Matrix!

Actual	Predictions	
	Won't Like	Will Like
Rated Poor	TN	FP
Rated High	FN	TP

KEY:

TN = true negative, FP = false positive
FN = false negative, TP = true positive

Row Sum = Total recommendations that could be made

Column Sum = Total recommendations that were made

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Deriving a Confusion Matrix

- Decide upon a rating threshold (T) to signify “likes”
 - E.g. A person likes a movie if they give it a rating ≥ 4
- Modify the test routine in order to keep 4 counts
 - TP (True Positive) ~ increment when predicted rating is $\geq T$ AND actual rating is $\geq T$
 - FP (False Positive) ~ increment when predicted rating is $\geq T$ AND actual rating is $< T$
 - TN (True Negative) ~ increment when predicted rating is $< T$ AND actual rating is $< T$
 - FN (False Negative) ~ increment when predicted rating is $< T$ AND actual rating is $\geq T$
- Increment the counts after each individual test (movie prediction) is made
- Display the counts as a confusion matrix at the end of the test

Coding the Confusion Matrix

```
#test recommendations for one user
testuserCM <- function(target, predthresh = 4, doprint=TRUE) {
  testitems = row.names(items[!is.na(items[,target]),])
  targetdata = items[testitems, target]
  names(targetdata) = testitems
  traindata = items[testitems, trainusernames]
  toterr = valid = TP = FP = TN = FN = 0
  for (item in testitems) {
    truerating = targetdata[item]
    targetdata[item] = NA
    sims = cor(targetdata, traindata, use="pairwise.complete.obs")
    sims = sims[,!is.na(sims)]
    prediction = weighted.mean(traindata[item, names(sims)], sims+1, na.rm=TRUE)
    if (!is.na(prediction)) {
      toterr = toterr + abs(prediction - truerating)
      valid = valid + 1
      if (prediction >= predthresh) {
        if (truerating >= predthresh) TP = TP + 1
        else FP = FP + 1
      }
      else if (truerating >= predthresh) FN = FN + 1
      else TN = TN + 1
    }
    targetdata[item] = truerating
  }
  if (doprint) cat(" avgerr=", toterr/valid, "#preds=", valid, "\n")
  return(c(toterr/valid, TP, FP, TN, FN))
}
```

← Add extra parameters

← Set all counts to zero

← Increment the TP, FP, TN, FN counts

← One way to return multiple values is to use a vector

Coding the Confusion Matrix

```
#test recommendations for all users, with confusion matrix
testallCM <- function() {
  TP = FP = TN = FN = toterr = cnt = 0
  for (user in testusernames) {
    res = testuserCM(user, doprint=FALSE)
    toterr = toterr + res[1]
    TP = TP + res[2]
    FP = FP + res[3]
    TN = TN + res[4]
    FN = FN + res[5]
    cnt = cnt + 1
    cat(user, res, "\n")
  }
  cat(sprintf("MAE= %0.4f TP=%d FP=%d TN=%d FN=%d\n", toterr/cnt, TP, FP, TN, FN))
}
```


Basic Item-Item CF Algorithm

- Precompute the similarities between all items (use Euclidean distance)

```
> itemsims = apply(items, 1, function(item)
+ apply(items, 1, function(x) 1/(1+sqrt(sum((x - item)^2,na.rm=TRUE)))))
>
> itemsims
```

	LadyInWater	SnakesOnPlane	JustMyLuck	Superman	Dupree	NightListener
LadyInWater	1.0000000	0.3483315	0.3483315	0.2402531	0.4721360	0.3761785
SnakesOnPlane	0.3483315	1.0000000	0.2553968	0.3090170	0.1876128	0.3266316
JustMyLuck	0.3483315	0.2553968	1.0000000	0.2079916	0.3761785	0.2708132
Superman	0.2402531	0.3090170	0.2079916	1.0000000	0.1846422	0.2742919
Dupree	0.4721360	0.1876128	0.3761785	0.1846422	1.0000000	0.2942981
NightListener	0.3761785	0.3266316	0.2708132	0.2742919	0.2942981	1.0000000

```
>
```

- The new getrecommendations() is:

```
getrecommendations2 <- function(username) {
  myRats = items[,username]
  wavrats = apply(itemsims, 1, function(simrow) weighted.mean(myRats, simrow, na.rm=TRUE))

  # remove items already rated by the user
  notseenitems = row.names(items[is.na(items[,username]),])
  t = wavrats[notseenitems] |
  sort(t[!is.na(t)] , decreasing = TRUE)[1:min(5,length(t))] # get top 5 items
}
```

Workshop Detailed Instructions

1. Execute and test the user-user CF code on the movielens data
 - Randomly select 20 test users and compute average MAE for their predicted ratings
 - Modify the functions testall() and testuser() in order to derive a confusion matrix using a “like” threshold of 4; compute the precision and recall
 - What is the impact of changing the “like” threshold? (e.g. making it 3) (what is the best trade-off between precision and recall?)
2. Repeat above using the similarity measures below (you will need to code the new measures). Which similarity measure gives the best results?
 - Euclidean Distance
 - Cosine Parallel
 - Pearson Coefficient - done already in (1)
3. Convert the functions testall() and testuser() to perform item-to-item CF. Do user-user and item-item CF give similar results?